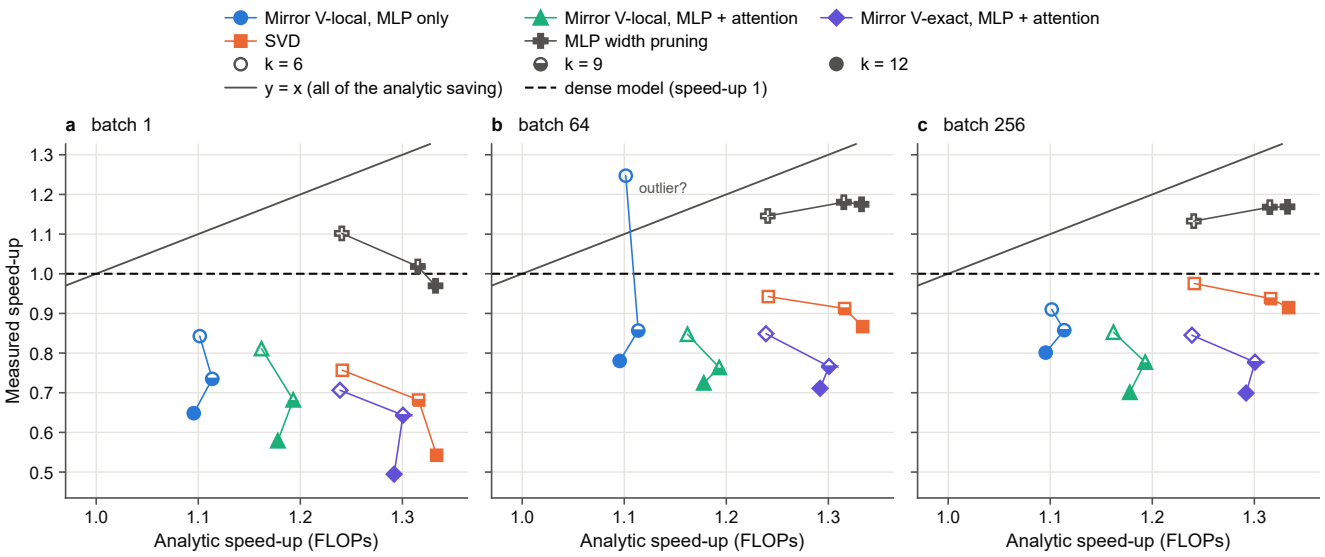
Mirror V-local, MLP only
Mirror V-local, MLP + attention
Mirror V-exact, MLP + attention
SVD
MLP width pruning
k = 6
k = 9
k = 12
y = x (all of the analytic saving)
dense model (speed-up 1)
a batch 1
b batch 64
c batch 256
outlier?
Measured speed-up
Analytic speed-up (FLOPs)
1.3
1.2
1.1
1.0
0.9
0.8
0.7
0.6
0.5